

Logistic Regression:

$$\hat{Y} = \frac{1}{1 + e^{-Z}}$$

$$Z = w.X + b$$

Y_hat --> predicted value

X --> Input Variable

w --> weight

b --> bias

Gradient Descent:

Gradient Descent is an optimization algorithm used for minimizing the loss function in various machine learning algorithms. It is used for updating the parameters of the learning model.

$$w = w - \alpha * dw$$

$$b = b - \alpha * db$$

Learning Rate:

Learning rate is a tuning parameter in an optimization algorithm that determines the step size at each iteration while moving toward a minimum of a loss function.

Derivatives:

$$dw = \frac{1}{m} * (\hat{Y} - Y). X$$

$$db = \frac{1}{m} * (\hat{Y} - Y)$$

Importing the Dependencies

```
In [1]: # importing numpy library
import numpy as np
```

Logistic Regression

```
In [2]: class LogisticRegression():

    # declaring learning rate & number of iterations (Hyperparameters)
    def __init__(self, learning_rate, no_of_iterations):

        self.learning_rate = learning_rate
        self.no_of_iterations = no_of_iterations

    # fit function to train the model with dataset
    def fit(self, X, Y):

        # number of data points in the dataset (number of rows) --> m
        # number of input features in the dataset (number of columns) --> n
        self.m, self.n = X.shape

        #initiating weight & bias value

        self.w = np.zeros(self.n)

        self.b = 0

        self.X = X

        self.Y = Y

        # implementing Gradient Descent for Optimization

        for i in range(self.no_of_iterations):
            self.update_weights()

    def update_weights(self):

        # Y_hat formula (sigmoid function)

        Y_hat = 1 / (1 + np.exp( - (self.X.dot(self.w) + self.b ) ))

        # derivatives

        dw = (1/self.m)*np.dot(self.X.T, (Y_hat - self.Y))

        db = (1/self.m)*np.sum(Y_hat - self.Y)

        # updating the weights & bias using gradient descent

        self.w = self.w - self.learning_rate * dw

        self.b = self.b - self.learning_rate * db
```

```
# Sigmoid Equation & Decision Boundary

def predict(self, X):

    Y_pred = 1 / (1 + np.exp( - (X.dot(self.w) + self.b ) ))
    Y_pred = np.where( Y_pred > 0.5, 1, 0)
    return Y_pred
```

Importing the Dependencies

```
In [3]: import pandas as pd
        from sklearn.preprocessing import StandardScaler
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import accuracy_score
```

Data Collection and Analysis

PIMA Diabetes Dataset

```
In [4]: # Loading the diabetes dataset to a pandas DataFrame
        diabetes_dataset = pd.read_csv(r'C:\Users\NITHISH\Desktop\DATA SCIENCE\NITISH-ML\
```

```
In [5]: # printing the first 5 rows of the dataset
        diabetes_dataset.head()
```

Out[5]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	
0	6	148	72	35	0	33.6	0.627	
1	1	85	66	29	0	26.6	0.351	
2	8	183	64	0	0	23.3	0.672	
3	1	89	66	23	94	28.1	0.167	
4	0	137	40	35	168	43.1	2.288	

```
In [6]: # number of rows and Columns in this dataset
        diabetes_dataset.shape
```

Out[6]: (768, 9)

```
In [7]: # getting the statistical measures of the data
diabetes_dataset.describe()
```

Out[7]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPe
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

```
In [8]: diabetes_dataset['Outcome'].value_counts()
```

Out[8]: 0 500
1 268
Name: Outcome, dtype: int64

0 --> Non-Diabetic

1 --> Diabetic

```
In [9]: diabetes_dataset.groupby('Outcome').mean()
```

Out[9]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
Outcome							
0	3.298000	109.980000	68.184000	19.664000	68.792000	30.304200	
1	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537	

```
In [10]: # separating the data and labels
features = diabetes_dataset.drop(columns = 'Outcome', axis=1)
target = diabetes_dataset['Outcome']
```

In [11]: `print(features)`

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age
0	0.627	50
1	0.351	31
2	0.672	32
3	0.167	21
4	2.288	33
..
763	0.171	63
764	0.340	27
765	0.245	30
766	0.349	47
767	0.315	23

[768 rows x 8 columns]

In [12]: `print(target)`

0	1
1	0
2	1
3	0
4	1
..	
763	0
764	0
765	0
766	1
767	0

Name: Outcome, Length: 768, dtype: int64

Data Standardization

In [13]: `scaler = StandardScaler()`

In [14]: `scaler.fit(features)`

Out[14]: `StandardScaler()`

In [15]: `standardized_data = scaler.transform(features)`

In [16]: `print(standardized_data)`

```
[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
   1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
  -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
  -0.10558415]
 ...
 [ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
  -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
   1.17073215]
 [-0.84488505 -0.8730192  0.04624525 ... -0.20212881 -0.47378505
  -0.87137393]]
```

In [17]: `features = standardized_data`
`target = diabetes_dataset['Outcome']`

In [18]: `print(features)`
`print(target)`

```
[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
   1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
  -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
  -0.10558415]
 ...
 [ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
  -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
   1.17073215]
 [-0.84488505 -0.8730192  0.04624525 ... -0.20212881 -0.47378505
  -0.87137393]]
```

0 1

1 0

2 1

3 0

4 1

..

763 0

764 0

765 0

766 1

767 0

Name: Outcome, Length: 768, dtype: int64

Train Test Split

In [19]: `X_train, X_test, Y_train, Y_test = train_test_split(features, target, test_size =`

```
In [20]: print(features.shape, X_train.shape, X_test.shape)
```

```
(768, 8) (614, 8) (154, 8)
```

Training the Model

```
In [21]: classifier = LogisticRegression(learning_rate=0.01, no_of_iterations=1000)
```

```
In [22]: #training the support vector Machine Classifier  
classifier.fit(X_train, Y_train)
```

Model Evaluation

Accuracy Score

```
In [23]: # accuracy score on the training data  
X_train_prediction = classifier.predict(X_train)  
training_data_accuracy = accuracy_score( Y_train, X_train_prediction)
```

```
In [24]: print('Accuracy score of the training data : ', training_data_accuracy)
```

```
Accuracy score of the training data :  0.7768729641693811
```

```
In [25]: # accuracy score on the test data  
X_test_prediction = classifier.predict(X_test)  
test_data_accuracy = accuracy_score( Y_test, X_test_prediction)
```

```
In [26]: print('Accuracy score of the test data : ', test_data_accuracy)
```

```
Accuracy score of the test data :  0.7662337662337663
```

Making a Predictive System

```
In [27]: input_data = (5,166,72,19,175,25.8,0.587,51)

# changing the input_data to numpy array
input_data_as_numpy_array = np.asarray(input_data)

# reshape the array as we are predicting for one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

# standardize the input data
std_data = scaler.transform(input_data_reshaped)
print(std_data)

prediction = classifier.predict(std_data)
print(prediction)

if (prediction[0] == 0):
    print('The person is not diabetic')
else:
    print('The person is diabetic')
```

[[0.3429808 1.41167241 0.14964075 -0.09637905 0.82661621 -0.78595734
 0.34768723 1.51108316]]

[1]
The person is diabetic

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\base.py:445: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names
warnings.warn(

In []: